The BookMarker Test Document

A small deterministic PDF used by the acceptance suite.

BookMarker wraps the marker engine behind a small HTTP service that converts one PDF at a time. This sample document gives the conversion pipeline plain text, a heading, and a simple table to work with. It is intentionally boring: the point is that the output is reproducible, so a direct CLI run and a service run can be compared line by line.

Each acceptance run converts this file through the service, converts the same file through the engine CLI with identical flags, and diffs the two markdown results. Differences should be limited to benign engine nondeterminism.

How to read this document

The first page carries running text. The second page carries a small bordered table. The third page carries a bullet list. Nothing here requires an image, so the sample stays small enough to travel inside the repository.

Chapter 1 - Tables and Figures

This chapter exists so the layout model has a simple bordered table to find, and so an LLM-enabled run has something the LLM path can annotate. The table below lists three fictional measurements with short notes.

Item	Value	Notes
alpha	12	first sample row
beta	34	second sample row
gamma	56	third sample row

Chapter 2 - Lists

- The service converts exactly one document at a time.
- A busy service refuses new work instead of queueing it.
- Uploaded files are deleted as soon as their job ends.
- An LLM-enabled run that cannot reach the LLM either fails or re-runs without it.

End of the BookMarker sample document.